



# RISC-V & HPC

Guided by :

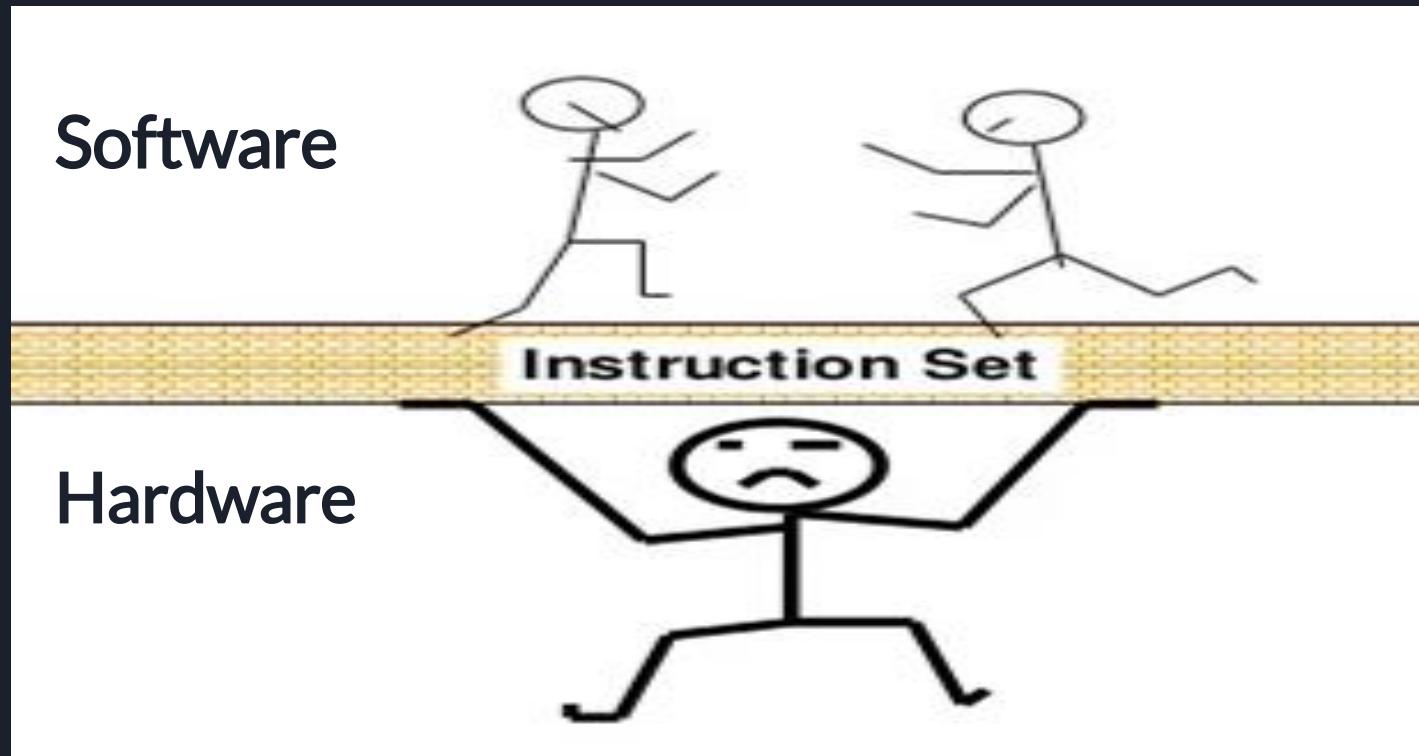
Mr. Surendra Billa Sir  
Mr. Rushikesh Jadhav Sir

Project P5:

Mr. Tejas Bhanudas Katkar  
Mr. Bhanu Prakash Agrawal  
Ms. Muskan Sharma  
Ms. Akanksha Abhay Ghorpade

PRN-230940130035  
PRN-230940130015  
PRN-230940130037  
PRN-230940130025

# Instruction Set Architecture





# Why RISC-V ?



# ARM vs RISC V: ISA Comparison

ISAs can be broadly categorized into two types:

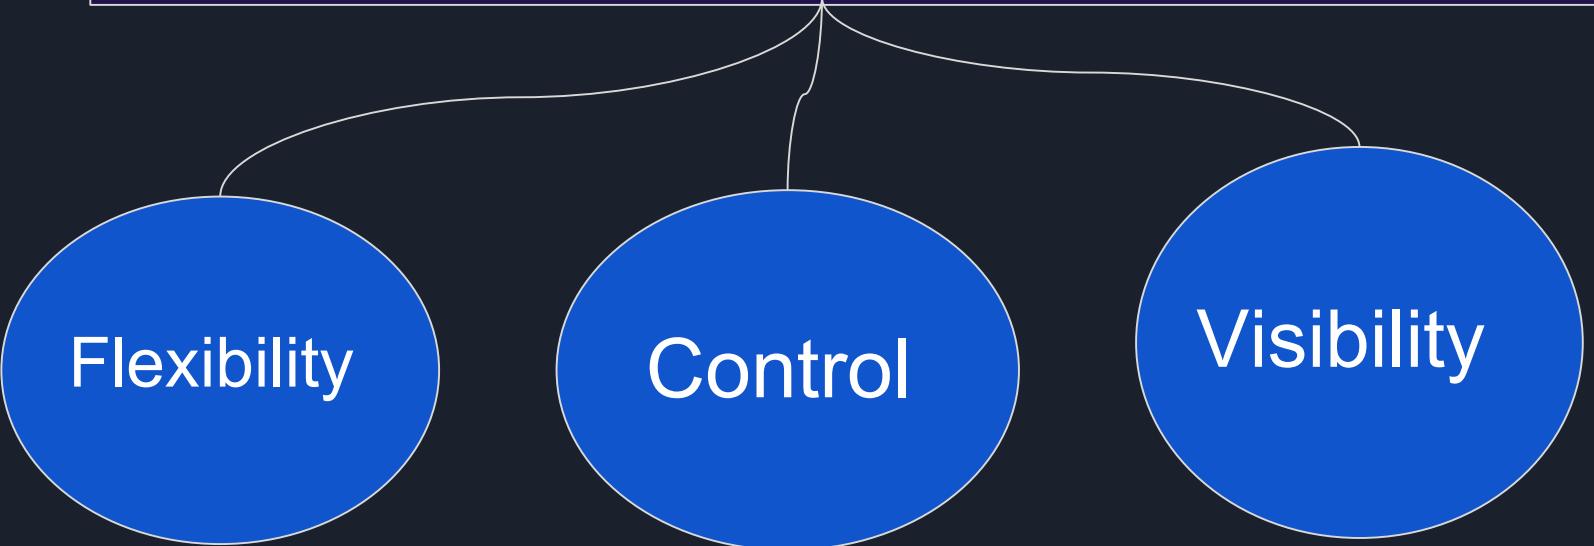
1→Open ISA

2→Closed ISA

- Closed ISAs, like ARM, are proprietary and tightly controlled by specific companies (Arm Holdings here), offering established reliability and compatibility but limiting customization.
- Open ISAs, exemplified by RISC-V, are community-driven and provide greater flexibility for customization, fostering innovation and adaptation to specific needs.



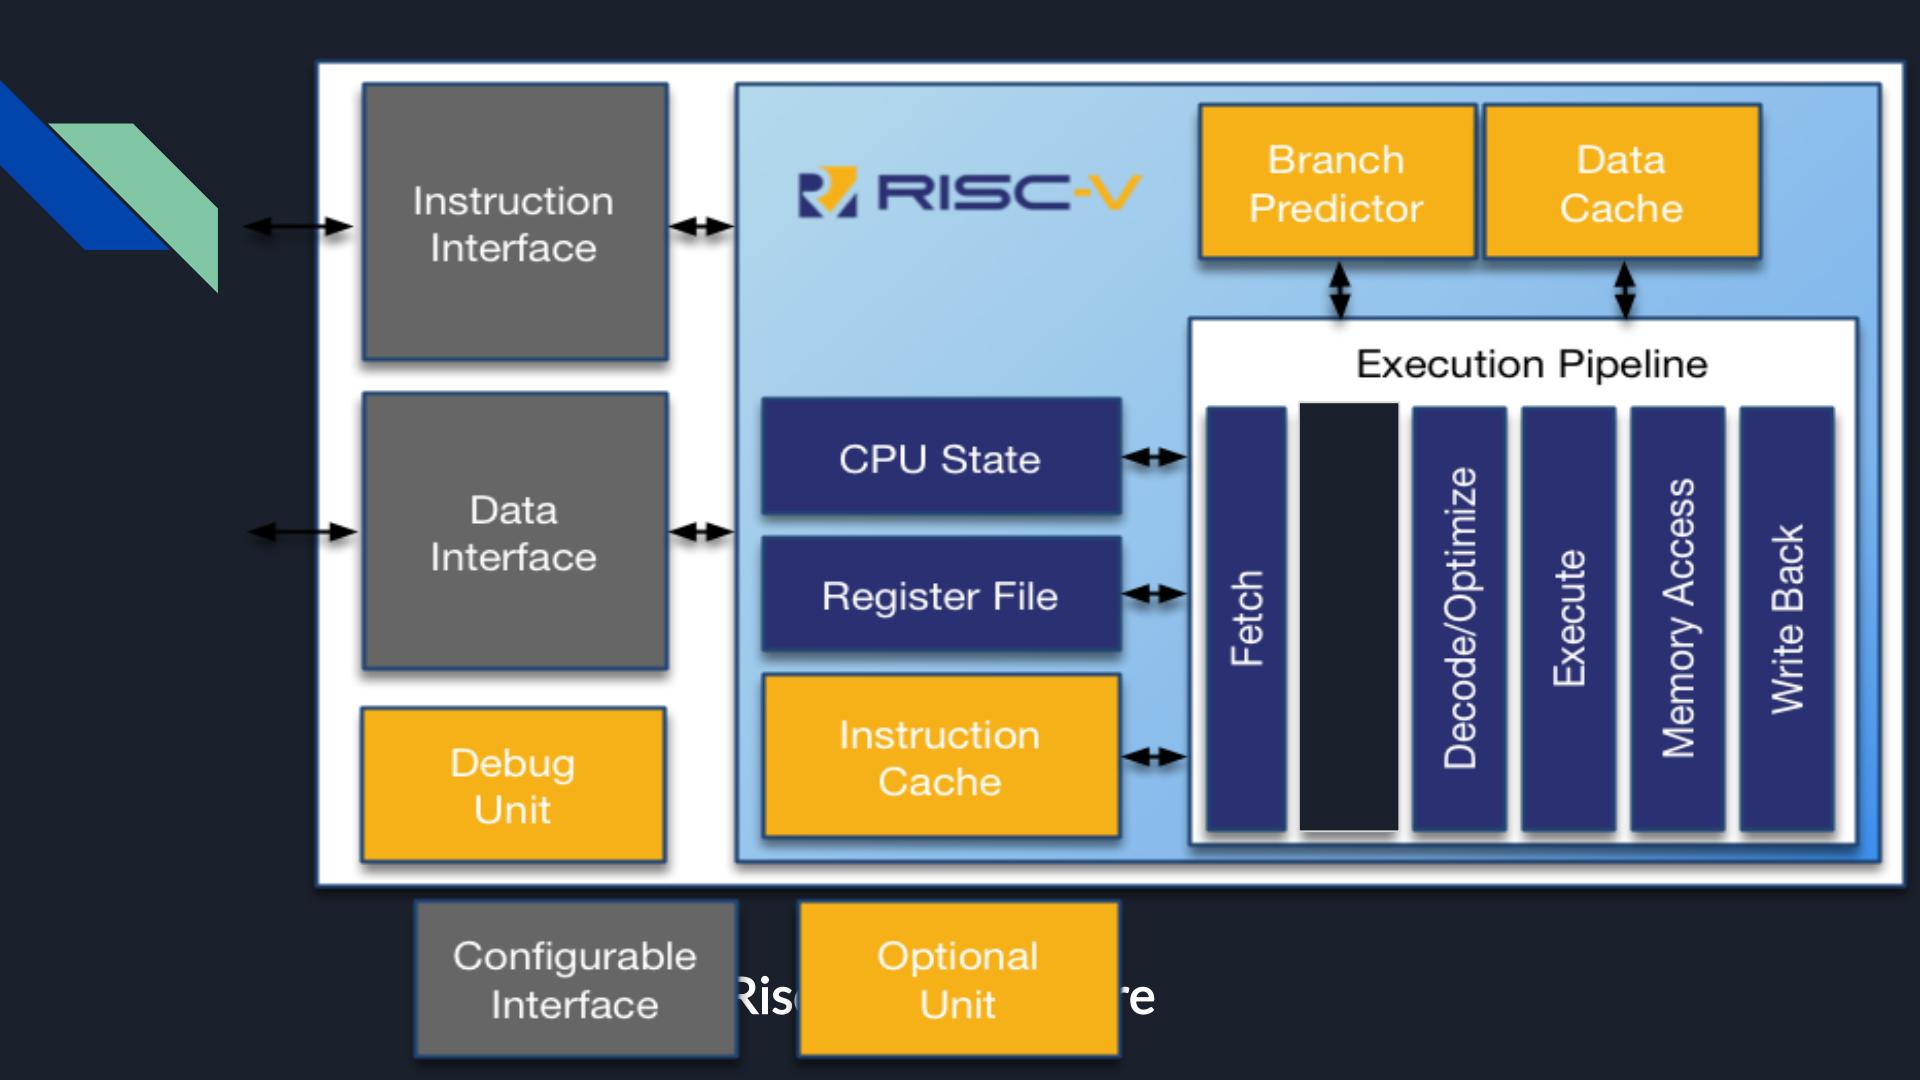
## 3 main advantages of RISC-V



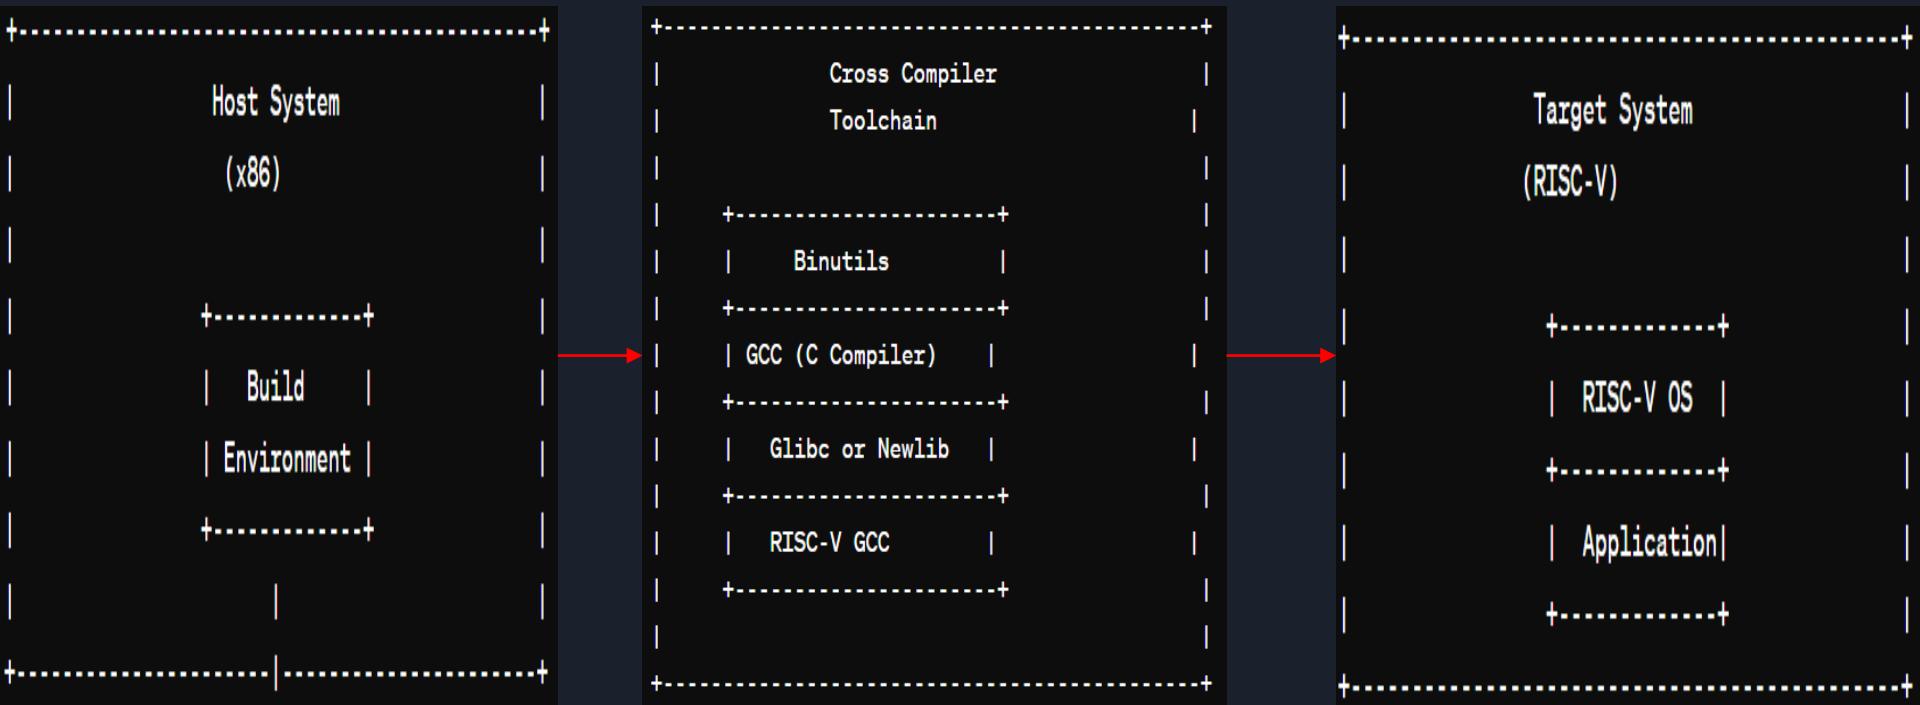
Flexibility

Control

Visibility



# Cross Compiler Toolchain (Flow)



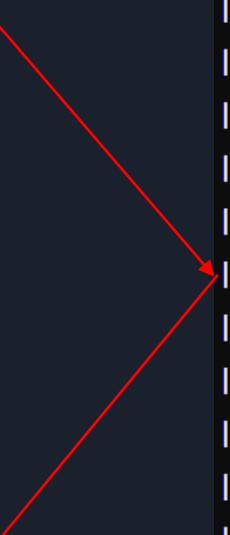


# Vector Processing in HPC

```
+-----+  
|       High-Performance Computing |  
|           (HPC) Application     |  
+-----+
```

```
+-----+  
|       Software Ecosystem        |  
| - Compilers, Libraries, Applications |  
| - Automatic Vectorization       |  
+-----+
```

```
+-----+  
|       RISC-V Vector Extension  |  
|           (RVV)                |  
+-----+  
| |       Vector Operations      | |  
| | - Vector Add, Multiply, Load/Store | |  
+-----+  
| |       Scalable Vector Lengths | |  
| | - Flexibility in Vector Size | |  
+-----+  
| |       Predication and Masking | |  
| | - Conditional Execution of Vectors | |  
+-----+  
| |  
+-----+
```





About RISC-V ▾ Membership ▾ RISC-V Exchange ▾ Technical ▾ News & Events ▾ Community ▾ Careers & Learning

Languages ▾ Tech Meetings Community Meetings Working Groups Portal Join



CINECA





# THEJAS32 SoC



**THEJAS32 SoC** is based on the **VEGA ET1031 processor**, which is a 32 bit single core in-order, 3-stage pipeline processor. This processor is based on the open source **RISC-V Instruction Set Architecture** and operates at a frequency of 100MHz. The SoC also includes peripherals like 256KB internal SRAM, Three UARTs, Four SPIs, Three TIMERS, Eight PWMs, Three I2C interfaces, 32 GPIOs etc. This SoC is designed and developed by Centre for Development of Advanced Computing (C-DAC) as part of the Digital India RISC-V (DIR-V) Program, by the Ministry of Electronics and Information Technology, Government of India.

# SIMD : Single instruction Multiple Data

How SIMD Operates ?

(a) Scalar Operation

$$\begin{array}{c}
 A_0 + B_0 = C_0 \\
 A_1 + B_1 = C_1 \\
 A_2 + B_2 = C_2 \\
 A_3 + B_3 = C_3
 \end{array}$$

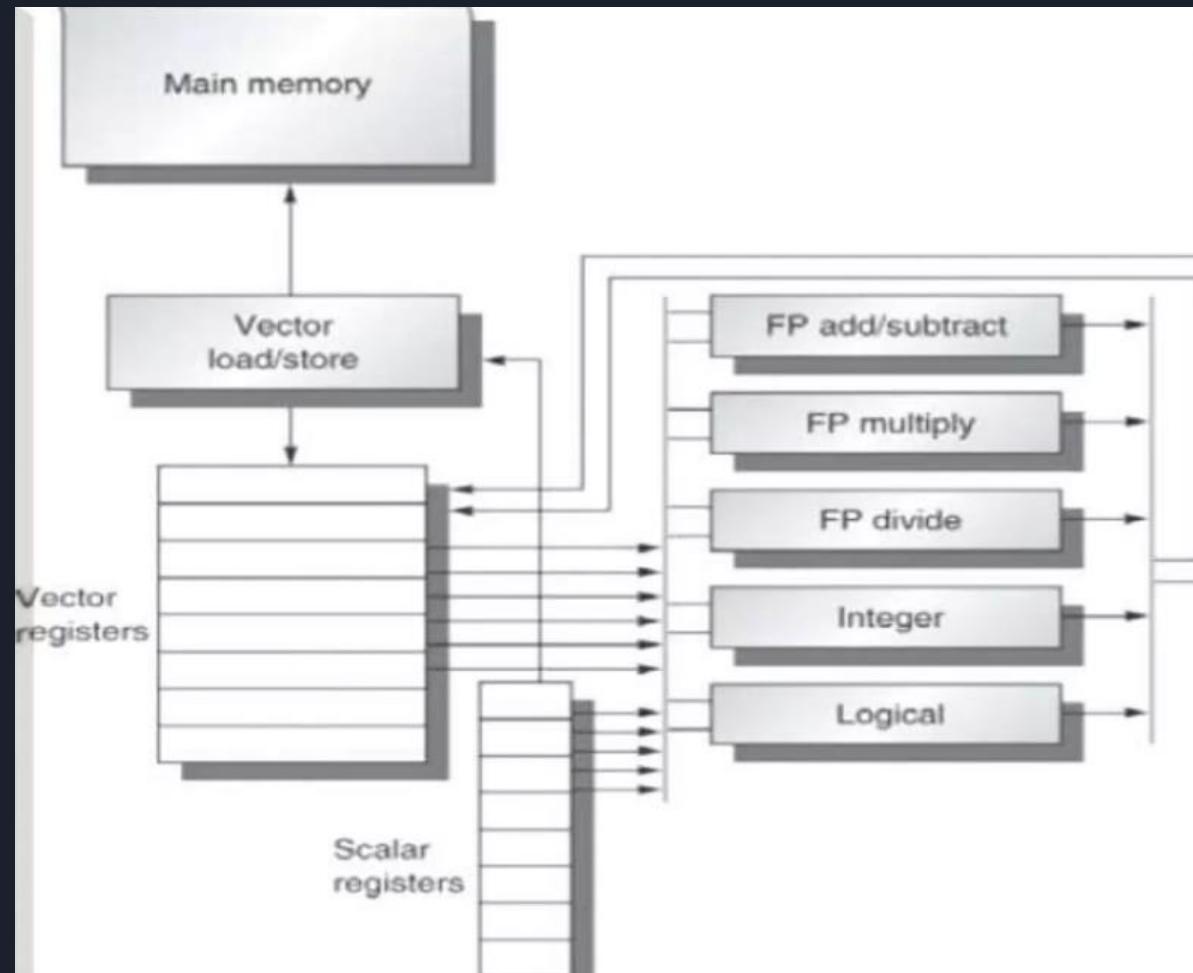
(b) SIMD Operation

$$\begin{array}{c}
 A_0 \\
 A_1 \\
 A_2 \\
 A_3
 \end{array}
 +
 \begin{array}{c}
 B_0 \\
 B_1 \\
 B_2 \\
 B_3
 \end{array}
 =
 \begin{array}{c}
 C_0 \\
 C_1 \\
 C_2 \\
 C_3
 \end{array}$$



# Vector Processing

- Vector Processing Architecture
  - Memory to Memory Architecture (Traditional)
  - Register to Register Architecture (Modern)
- Components of vector processors
  - Vector Registers
  - Vector functional Units
  - Vector load-store Units
  - Scalar Registers
- Advantages of Vector processing
  - Quick fetch and decode of single instruction for multiple operations
  - Easier Addressing of main memory
  - Elimination of memory Wastage
  - Simplification of control hazards
  - Reduced code size



# RISC-V V/S RISC-V V

## Risc-v

- Base instruction Set Architecture
- Focuses in Scalar Operations
- Handles individually data elements
- Performance is less as compared to RVV
- Easy to understand
- Widely Adopted

## Risc-v v

- Optional extension adds capabilities of SIMD
- Focuses on Vector Operations
- Handles multiple data elements
- Designed to work on large data sets
- High Performance
- RVV support is growing but not universally implemented on all Risc-v

# RISC-V VECTOR EXTENSION

**Unleashing Parallelism in RISC-V  
Architecture**



# RISC-V EXTENSIONS

- RISC-V is augmented via the concept of extensions
  - Extensions can add new instructions and CPU state
- Base ISA is called I (for Integer)
  - RV32I
  - RV64I ( $XLEN=64$ , adds a few arithmetic instructions to improve 32-bit integer arithmetic)
- Common Standard Extensions in a RISC-V 64-bit Linux capable core
  - M. Integer multiplication and division (mul, div, rem, ...)
  - A. Atomic instructions (load reserve + store conditional, atomic read-modify-write)
  - F. Single-Precision Floating-Point (IEEE 754 Binary32)
  - D. Double-Precision Floating-Point (IEEE 754 Binary64)
  - C. Compressed Instructions (16-bit encodings for common I/F/D instructions)

IMAFD = G

# INTRODUCTION

- **What it is:** A set of instructions for RISC-V processors that enable them to perform vector operations.
- **What it does:** Allows processing multiple data elements at once, leading to faster performance for specific tasks.
- **When it's useful:** For tasks involving large amounts of data and repetitive calculations, such as:
  - Signal processing
  - Image and video processing
  - Machine learning
  - Scientific computing

# OVERVIEW OF ITS FEATURES AND CAPABILITIES

- 1. SIMD (Single Instruction, Multiple Data):** Allows a single instruction to perform the same operation on multiple data elements simultaneously, exploiting data-level parallelism.
- 2. Enhanced Performance:** By processing multiple data elements in parallel, the Vector Extension significantly improves throughput and efficiency for tasks such as multimedia processing, scientific computing, and machine learning.
- 3. Customizable Vector Length:** Supports vectors of variable lengths, enabling optimization for different applications and target platforms.
- 4. Vector Registers:** Dedicated vector registers hold vector data, facilitating efficient manipulation and computation.
- 5. Vector Memory Accesses:** Supports efficient vector memory accesses, reducing overhead in loading and storing vector data.
- 6. Predication Support:** Enables conditional execution based on vector elements, enhancing flexibility in control flow.

# RISC-V RV64GV TOOLCHAIN

```
muskan@cdackoji:~$ ls
lj-os  muskan  packages  packages1  qemu  riscv  self
muskan@cdackoji:~$ cd muskan/opt/bin/
muskan@cdackoji:~/muskan/opt/bin$ ls
riscv64-unknown-elf-addr2line      riscv64-unknown-elf-lto-dump
riscv64-unknown-elf-ar             riscv64-unknown-elf-nm
riscv64-unknown-elf-as             riscv64-unknown-elf-objcopy
riscv64-unknown-elf-c++            riscv64-unknown-elf-objdump
riscv64-unknown-elf-c++filt        riscv64-unknown-elf-ranlib
riscv64-unknown-elf-cpp            riscv64-unknown-elf-readelf
riscv64-unknown-elf-elfedit        riscv64-unknown-elf-run
riscv64-unknown-elf-g++            riscv64-unknown-elf-size
riscv64-unknown-elf-gcc            riscv64-unknown-elf-strings
riscv64-unknown-elf-gcc-13.2.0     riscv64-unknown-elf-strip
riscv64-unknown-elf-gcc-ar         riscv64-unknown-linux-gnu-addr2line
riscv64-unknown-elf-gcc-nm         riscv64-unknown-linux-gnu-ar
riscv64-unknown-elf-gcc-ranlib     riscv64-unknown-linux-gnu-as
riscv64-unknown-elf-gcov           riscv64-unknown-linux-gnu-c++
riscv64-unknown-elf-gcov-dump      riscv64-unknown-linux-gnu-c++filt
riscv64-unknown-elf-gcov-tool      riscv64-unknown-linux-gnu-cpp
riscv64-unknown-elf-gdb            riscv64-unknown-linux-gnu-elfedit
riscv64-unknown-elf-gdb-add-index   riscv64-unknown-linux-gnu-g++
riscv64-unknown-elf-gprof          riscv64-unknown-linux-gnu-gcc
riscv64-unknown-elf-ld              riscv64-unknown-linux-gnu-gcc-13.2.0
riscv64-unknown-elf-ld.bfd         riscv64-unknown-linux-gnu-gcc-ar
riscv64-unknown-elf-nm             riscv64-unknown-linux-gnu-gcov
riscv64-unknown-elf-objcopy        riscv64-unknown-linux-gnu-gcov-dump
riscv64-unknown-elf-objdump        riscv64-unknown-linux-gnu-gcov-tool
riscv64-unknown-elf-readelf        riscv64-unknown-linux-gnu-gdb
riscv64-unknown-elf-run            riscv64-unknown-linux-gnu-gdb-add-index
riscv64-unknown-elf-size           riscv64-unknown-linux-gnu-gfortran
riscv64-unknown-elf-strings        riscv64-unknown-linux-gnu-gprof
riscv64-unknown-elf-strip          riscv64-unknown-linux-gnu-ld
riscv64-unknown-linux-gnu-addr2line riscv64-unknown-linux-gnu-ld.bfd
riscv64-unknown-linux-gnu-ar       riscv64-unknown-linux-gnu-lto-dump
riscv64-unknown-linux-gnu-as       riscv64-unknown-linux-gnu-nm
riscv64-unknown-linux-gnu-c++      riscv64-unknown-linux-gnu-objcopy
riscv64-unknown-linux-gnu-c++filt  riscv64-unknown-linux-gnu-objdump
riscv64-unknown-linux-gnu-cpp      riscv64-unknown-linux-gnu-ranlib
riscv64-unknown-linux-gnu-elfedit   riscv64-unknown-linux-gnu-readelf
riscv64-unknown-linux-gnu-g++       riscv64-unknown-linux-gnu-run
riscv64-unknown-linux-gnu-gcc       riscv64-unknown-linux-gnu-size
riscv64-unknown-linux-gnu-gcc-13.2.0 riscv64-unknown-linux-gnu-strings
riscv64-unknown-linux-gnu-gcc-ar    riscv64-unknown-linux-gnu-strip
```

# RISC-V VECTOR INSTRUCTION

```
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <sys/time.h>

#define MAXELEM      1000

uint64_t get_time_us() {
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return (tv.tv_sec * 1000000) + tv.tv_usec;
}

extern void vvaddint32(size_t n, const int *x, const int *y, int *z);

int main()
{
    int x[MAXELEM] = {0};
    int y[MAXELEM] = {0};
    int z[MAXELEM] = {0};

    srand(time(NULL));

    for(int i=0; i < MAXELEM; i++)
    {
        x[i] = rand() % 1000;
        y[i] = rand() % 1000;
    }

    uint64_t begin = get_time_us();
    vvaddint32(MAXELEM, x, y, z);
    uint64_t elapsed = get_time_us() - begin;

    printf("\nTime taken for vector addition : %lu us\n", elapsed);

}
```

```
.text
.balign 4
.global vvaddint32
# vector-vector add routine of 32-bit integers
# void vvaddint32(size_t n, const int*x, const int*y, int*z)
# { for (size_t i=0; i<n; i++) { z[i]=x[i]+y[i]; } }
#
# a0 = n, a1 = x, a2 = y, a3 = z
# Non-vector instructions are indented
vvaddint32:
    vsetvli t0, a0, e32, ta, ma # Set vector length based on 32-bit vectors
    vle32.v v0, (a1)           # Get first vector
    sub a0, a0, t0              # Decrement number done
    slli t0, t0, 2              # Multiply number done by 4 bytes
    add a1, a1, t0              # Bump pointer
    vle32.v v1, (a2)           # Get second vector
    add a2, a2, t0              # Bump pointer
    vadd.vv v2, v0, v1          # Sum vectors
    vse32.v v2, (a3)           # Store result
    add a3, a3, t0              # Bump pointer
    bnez a0, vvaddint32         # Loop back
    ret                         # Finished
```

# RISC-V SCALAR INSTRUCTION

```
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <sys/time.h>

#define MAXELEM      1000

uint64_t get_time_us() {
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return (tv.tv_sec * 1000000) + tv.tv_usec;
}

extern void vvaddint32(size_t n, const int *x, const int *y, int *z);

int main()
{
    int x[MAXELEM] = {0};
    int y[MAXELEM] = {0};
    int z[MAXELEM] = {0};

    srand(time(NULL));

    for(int i=0; i < MAXELEM; i++)
    {
        x[i] = rand() % 1000;
        y[i] = rand() % 1000;
    }

    uint64_t begin = get_time_us();
    for(int i=0; i< MAXELEM; i++)
    {
        z[i] = x[i] + y[i];
    }
    uint64_t elapsed = get_time_us() - begin;

    printf("\nTime taken for non-vector addition : %lu us\n", elapsed);
}
```

# RISC-V MAKEFILE INSTRUCTION

```
CC = riscv64-unknown-elf-gcc  
|  
bin:  
    $(CC) -c vvadd.s  
    $(CC) -c main.c  
    $(CC) main.o vvadd.o -o main  
    $(CC) novect.c -o novect  
  
run:  
    @-qemu-riscv64-static -cpu rv64,v=true main  
    @-qemu-riscv64-static -cpu rv64,v=true novect
```

# EXPLORING RISC-V VECTOR EXTENSION WITH QEMU

```
muskan@cdackoji:~/self$ ls
main main.c main.o Makefile novect novect.c vvadd.o vvadd.s
muskan@cdackoji:~/self$ make
riscv64-unknown-elf-gcc -c vvadd.s
riscv64-unknown-elf-gcc -c main.c
riscv64-unknown-elf-gcc main.o vvadd.o -o main
riscv64-unknown-elf-gcc novect.c -o novect
muskan@cdackoji:~/self$ make run
```

Time taken for vector addition : 339 us

Time taken for non-vector addition : 107 us

```
muskan@cdackoji:~/self$ |
```

# X\_86 ARCHITECTURE

## Vector Instruction

```
#include <immintrin.h>
#include <stdio.h>
#include <sys/time.h>
#include <stdint.h>

#define n 100000
uint64_t get_time_us() {
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return (tv.tv_sec * 1000000) + tv.tv_usec;
}

int main() {
    // Initialize vectors with constant values
    __m256d veca = _mm256_set1_pd(6.0);
    __m256d vecb = _mm256_set1_pd(2.0);
    __m256d vecc = _mm256_set1_pd(7.0);

    // Initialize result array to store 1000 result values
    double result[n];

    uint64_t begin = get_time_us();
    // Perform the operation for 1000 entries
    for (int i = 0; i < n; i += 4) {
        __m256d result_vec = _mm256_fmaddsub_pd(veca, vecb, vecc);
        _mm256_storeu_pd(&result[i], result_vec);
    }

    uint64_t elapsed = get_time_us() - begin;

    // Display a few elements of the result array
    printf("Results for first few elements:\n");
    for (int i = 0; i < n; i++) {
        printf("%lf ", result[i]);
    }
    printf("\n");
    printf("\nTime taken for vector addition : %lu us\n", elapsed);

    return 0;
}
```

## Scalar Instruction

```
#include <stdio.h>
#include <sys/time.h>
#include <stdint.h>

#define NUM_ENTRIES 100000

double get_time_us() {
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return (tv.tv_sec * 1000000) + tv.tv_usec;
}

void fmaddsub(double *a, double *b, double *c, double *result, int n) {
    for (int i = 0; i < n; i++) {
        result[i] = a[i] * b[i] + ((i % 2 == 0) ? c[i] : -c[i]);
    }
}

int main() {
    double veca[NUM_ENTRIES];
    double vecb[NUM_ENTRIES];
    double vecc[NUM_ENTRIES];
    double result[NUM_ENTRIES];

    // Initialize vectors
    for (int i = 0; i < NUM_ENTRIES; i++) {
        veca[i] = 6.0;
        vecb[i] = 2.0;
        vecc[i] = 7.0;
    }

    // Perform the operation
    uint64_t begin = get_time_us();
    fmaddsub(veca, vecb, vecc, result, NUM_ENTRIES);
    uint64_t elapsed = get_time_us() - begin;

    // Display the result
    printf("Result for the first %d entries:\n", NUM_ENTRIES);
    for (int i = 0; i < NUM_ENTRIES; i++) {
        printf("%lf ", result[i]);
    }
    printf("\n");
    printf("\nTime taken for normal c operation: %lu us\n", elapsed);

    return 0;
}
```

# EXPLORING VECTOR EXTENSION WITH X\_86 ARCHITECTURE

```
muskan@cdackoji:~/muskan/x86_example$ ls  
c_code    old      vector_code111  x86_vector_ext.c  
c_code111 vector_code  x86_c_code.c  
muskan@cdackoji:~/muskan/x86_example  ./vector_code  
Results for first few elements:  
5.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.  
.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.0  
00000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.00  
0000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.0000
```

Time taken for vector addition : 3 us

प्रकाशितप्राचीन / प्रकाश / ४०० एवंप्रते

```
muskan@cdackoji:~/muskan/x86_example$ ls  
c_code    old      vector_code11 x86_vector_ext.  
c_code11  vector_code  x86_c_code.c  
muskan@cdackoji:~/muskan/x86_example$ ./c_code
```

.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000  
000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000  
0000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000  
000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000  
0 19.000000 5.000000 19.000000 5.000000 19.000000 5.000000 19.000000  
5.000000 19.000000 5.000000 19.000000 5.000000

Time taken for normal c operation: 13 us

Digitized by srujanika@gmail.com

# CHALLENGES AND LIMITATIONS

- 1. Programming Complexity:** Utilizing vector instructions effectively requires specialized programming techniques such as loop unrolling, data alignment, and explicit vectorization, which can increase code complexity and development time.
- 2. Data Dependency:** Dependencies among vector elements can limit parallelism and lead to inefficient utilization of vector resources, especially in algorithms with irregular data access patterns or conditional branching.
- 3. Memory Bandwidth:** Vectorized operations can impose significant demands on memory bandwidth, especially for large vector lengths, leading to potential bottlenecks and inefficiencies in memory-bound applications.
- 4. Hardware Complexity:** Implementing a high-performance vector unit entails additional hardware resources and complexity, including wider data paths, vector registers, and efficient memory access mechanisms, which may increase design complexity and power consumption.
- 5. Software Ecosystem:** The availability of software tools, libraries, and compilers optimized for the Vector Extension may be limited compared to established architectures, posing challenges for developers in leveraging its full potential.

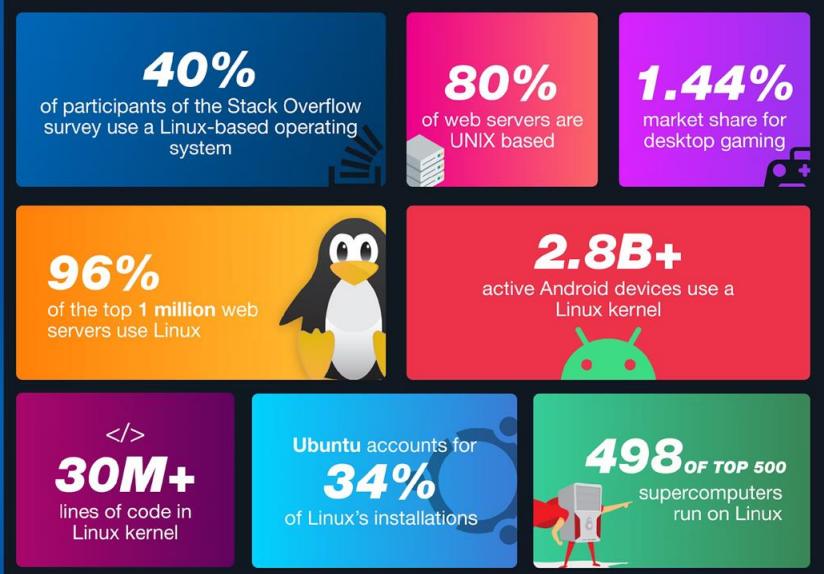
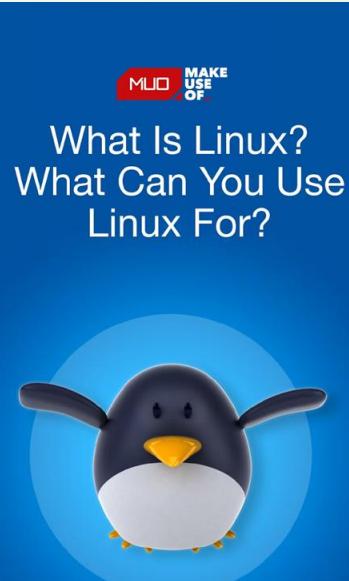
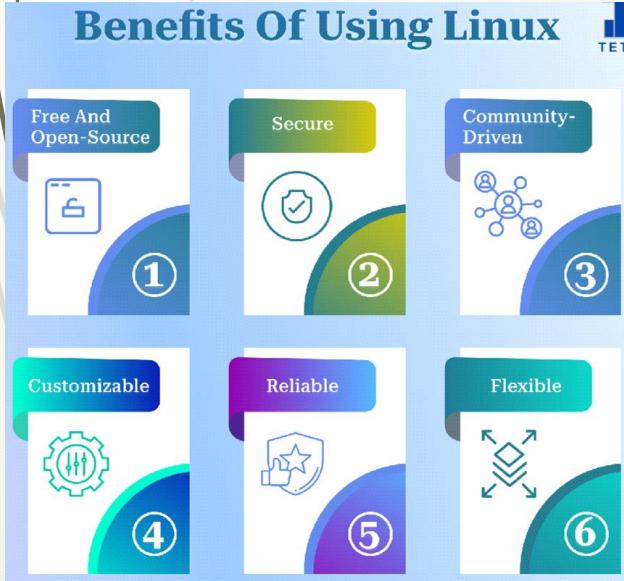
# AREAS FOR IMPROVEMENT AND ONGOING RESEARCH

- 1. Compiler Optimization:** Continued research is needed to develop advanced compiler techniques for automatic vectorization, loop restructuring, and optimization of memory access patterns to extract maximum parallelism from applications.
- 2. Vectorization Libraries:** Further development of vectorization libraries and frameworks tailored for the Vector Extension can simplify programming and accelerate adoption by providing high-level abstractions and optimized routines for common tasks.
- 3. Memory Hierarchy:** Research into efficient data prefetching, caching strategies, and memory hierarchy designs is crucial to mitigate the impact of memory bandwidth limitations and improve overall system performance.
- 4. Scalability and Heterogeneity:** Investigating techniques for scalable vector processing across multiple cores and heterogeneous architectures, including integration with multi-core processors and accelerators, can enhance performance and flexibility in diverse computing environments.
- 5. Energy Efficiency:** Optimizing vector architectures for energy efficiency through techniques such as dynamic voltage and frequency scaling, power gating, and instruction-level energy reduction is essential for enabling high-performance, energy-efficient computing systems.

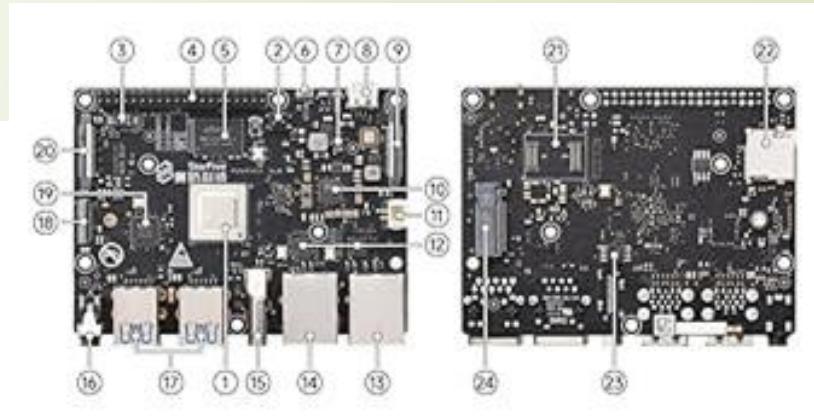
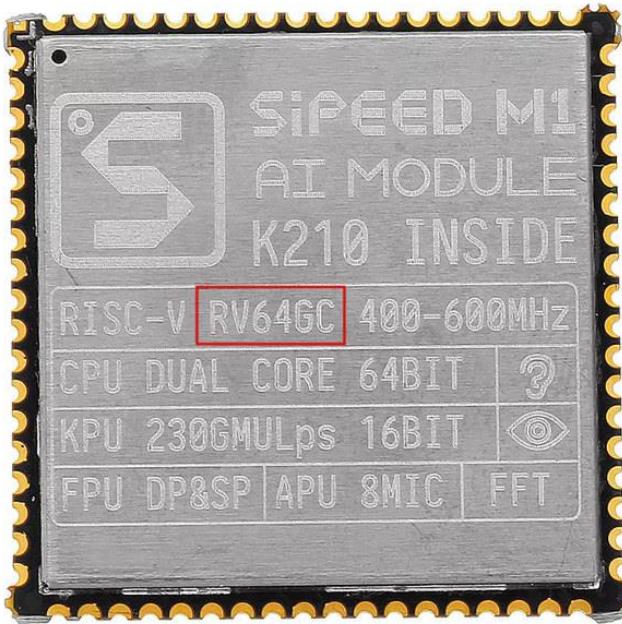
# DIY LINUX – What?



# DIY LINUX – Why?



# DIY LINUX – Where? – RVV (RISC-V Vector) riscv64-gv (RV64GV)



```
● ● jaufranc@FX8350: /media/hdd/edev/sandbox/riscv-linux/release
ucbvx login: root
Password:
root@ucbvx:~# uname -a
Linux ucbvx 4.14.0-00032-gd10799b22913-dirty #79 Sun Dec 17 10:41:31 NZDT 2017
riscv64 GNU/Linux
root@ucbvx:~# cat /proc/cpuinfo
hart : 0
tsa : rv64imafdsu
mmu : sv48

root@ucbvx:~# df -h
Filesystem      Size   Used  Available Use% Mounted on
/dev/root       58.0M  44.9M     8.6M  84% /dev
devtmpfs        59.8M     0    59.8M  0% /dev
none            59.9M     0    59.9M  0% /tmp
none            59.9M     0    59.9M  0% /var/tmp
root@ucbvx:~#
```

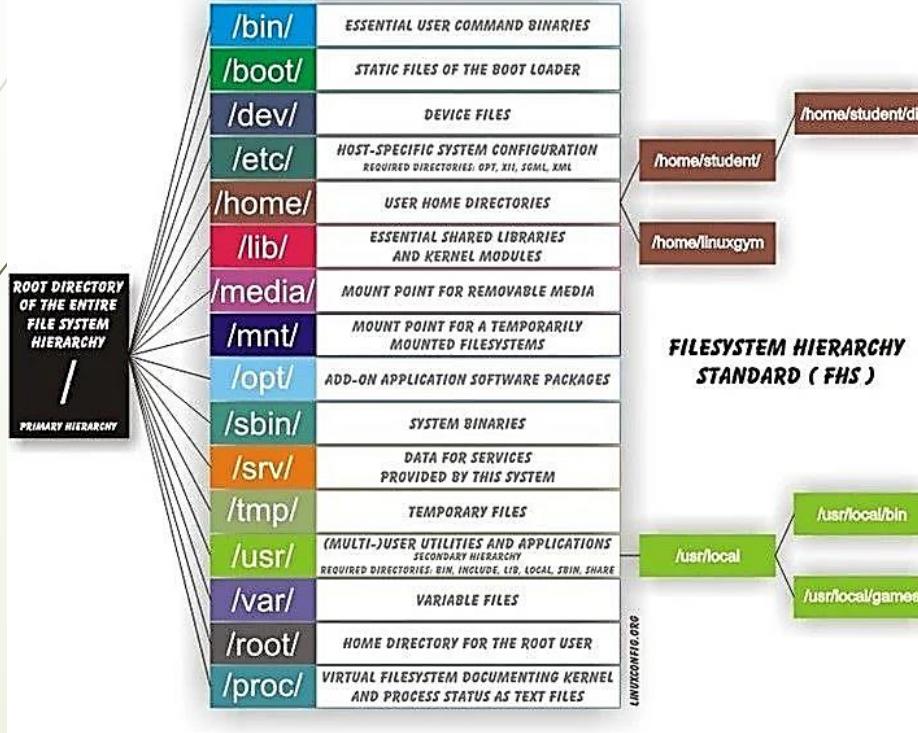


# DIY LINUX – How?

- 1) Configuring the Environment
- 2) Making cross Toolchain

# DIY LINUX – How?

## 1) Configuring the Environment



```
bhanu@fedora:/lj-os$ tree -d -L 1
```

```
.
```

```
bin
```

```
boot
```

```
cross-tools
```

```
dev
```

```
etc
```

```
home
```

```
lib
```

```
lib64
```

```
media
```

```
mnt
```

```
opt
```

```
proc
```

```
root
```

```
sbin
```

```
srv
```

```
sys
```

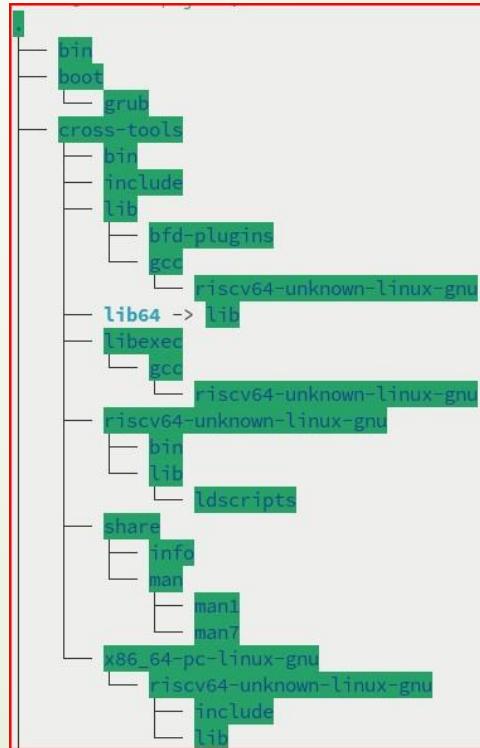
```
tmp
```

```
usr
```

```
var
```

# DIY LINUX – How?

## 1) Configuring the Environment 1) bin

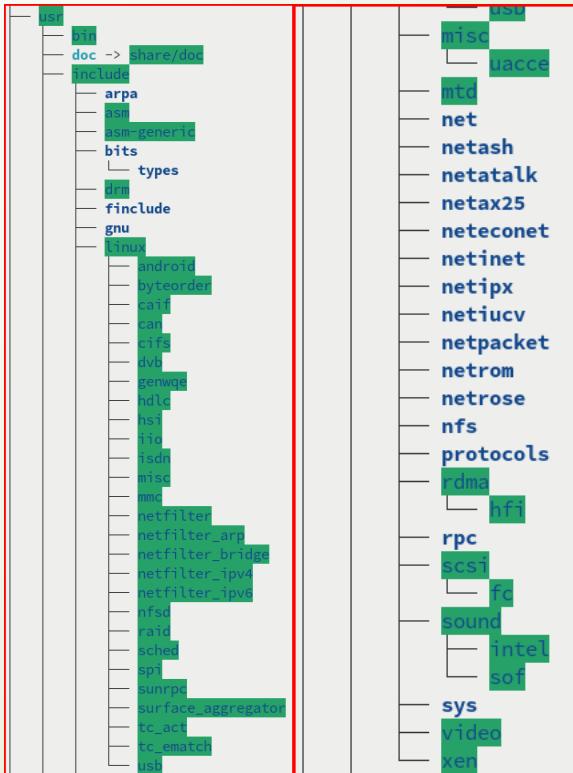


```
bhanu@fedora:/lj-os$ tree -d -L 1
.:
bin
boot
cross-tools
dev
etc
home
lib
lib64
media
mnt
opt
proc
root
sbin
srv
sys
tmp
usr
var
```

# DIY LINUX – How?

## 1) Configuring the Environment

### 1) usr

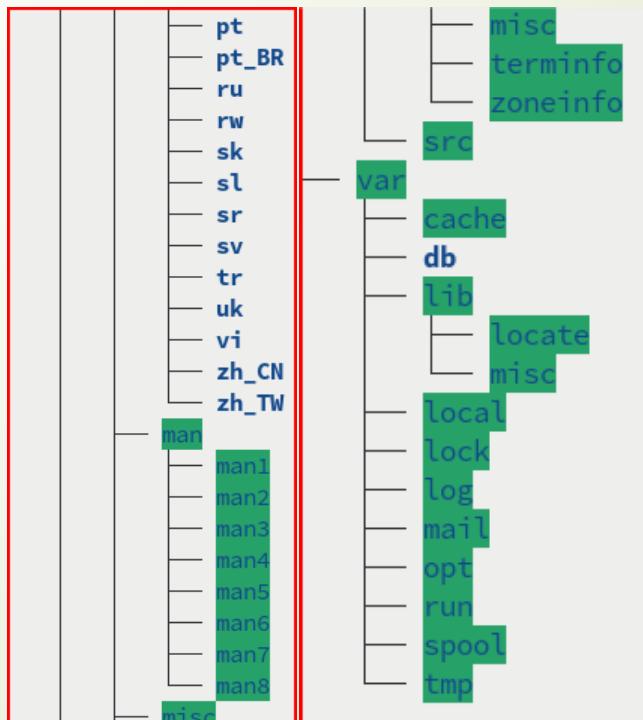


```
bhanu@fedora:/lj-os$ tree -d -L 1
.
├── bin
├── boot
├── cross-tools
├── dev
├── etc
├── home
├── lib
├── lib64
├── media
├── mnt
├── opt
├── proc
├── root
├── sbin
├── srv
├── sys
├── tmp
├── usr
└── var
```

# DIY LINUX – How?

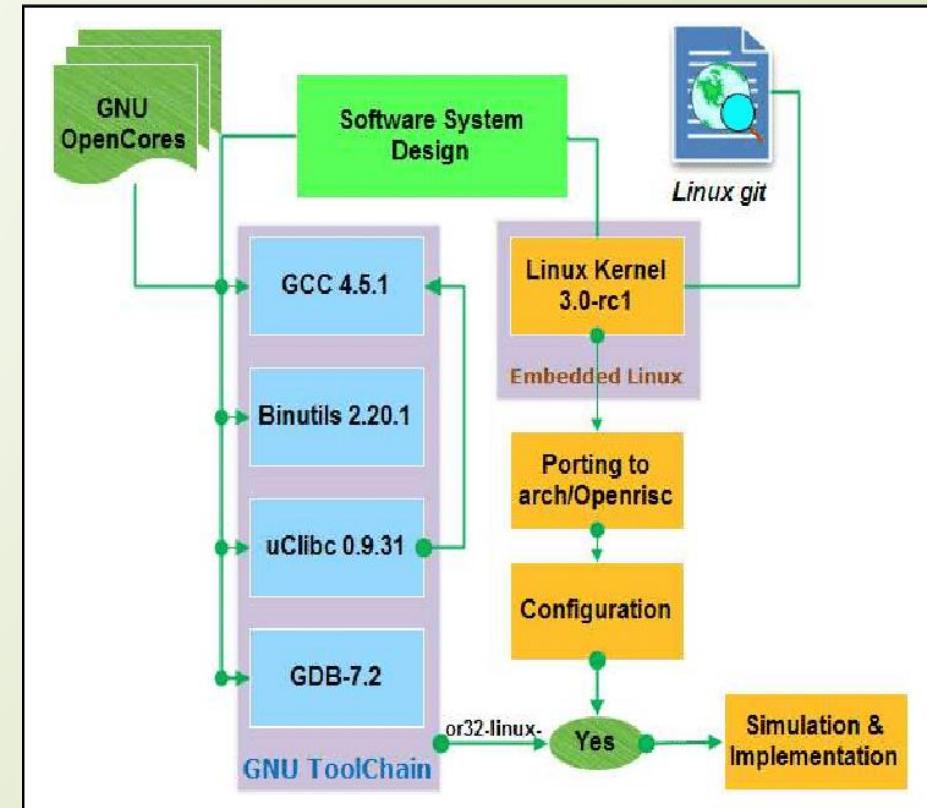
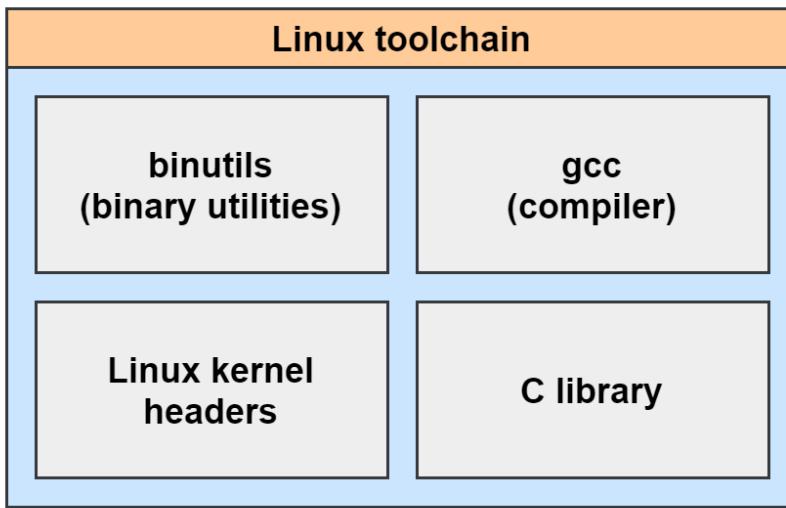
## 1) Configuring the Environment

### 1) var



```
bhanu@fedora:/lj-os$ tree -d -L 1
.
├── bin
├── boot
├── cross-tools
├── dev
├── etc
├── home
├── lib
├── lib64
├── media
├── mnt
├── opt
├── proc
├── root
├── sbin
├── srv
├── sys
├── tmp
└── usr
    └── var
```

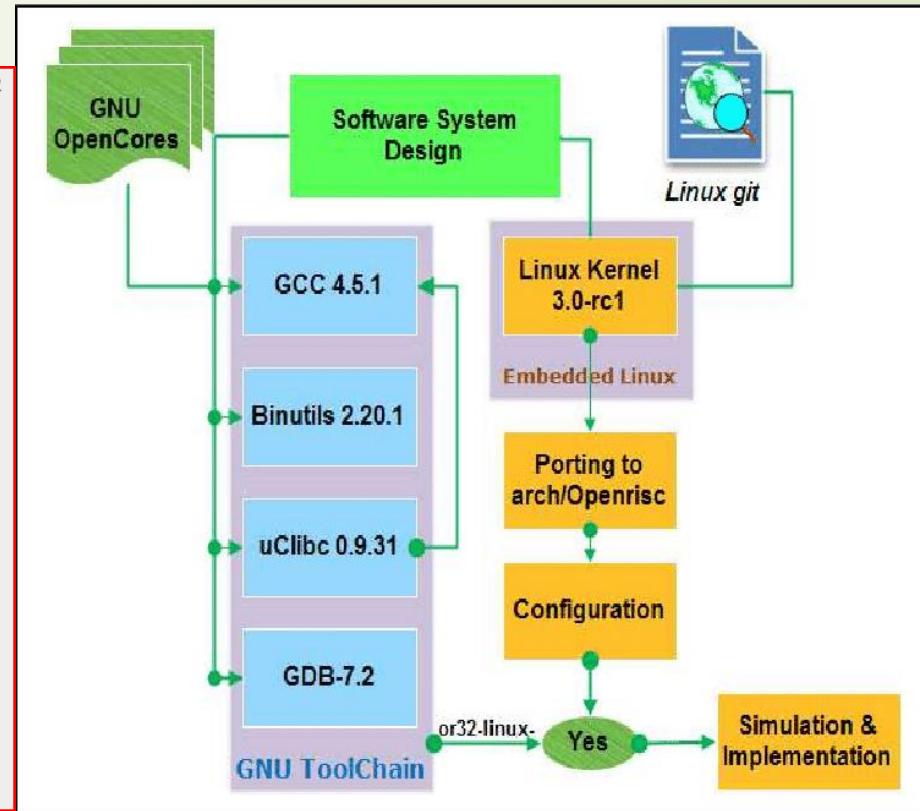
# DIY LINUX – How? – Toolchain



# DIY LINUX – How? –

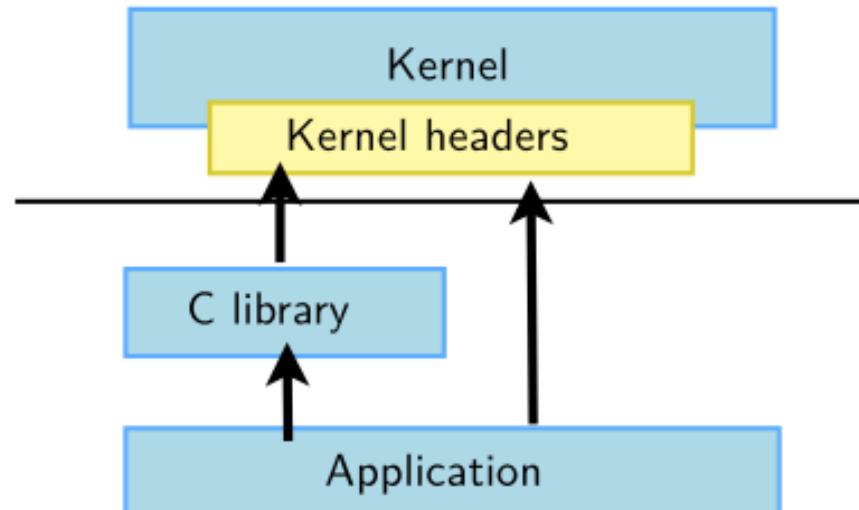
## Toolchain

```
bhanu@fedora:~/Documents/tars_glibc/tars$ tree -d -L 2
.
├── binutils
│   ├── binutils-2.39
│   └── binutils-build
├── busybox
├── clfs
├── gcc
│   ├── gcc-13.0.1-20230401
│   ├── gcc-build
│   └── gcc-static
├── glibc
│   ├── glibc-2.37
│   └── glibc-build
├── gmp
│   └── gmp-6.2.1
├── isl
│   └── isl-0.24
├── linux
│   └── linux-6.7.3
├── mpc
│   └── mpc-0.34
├── mpfr
│   └── mpfr-4.1.1
└── zlib
```



# DIY LINUX – How? –

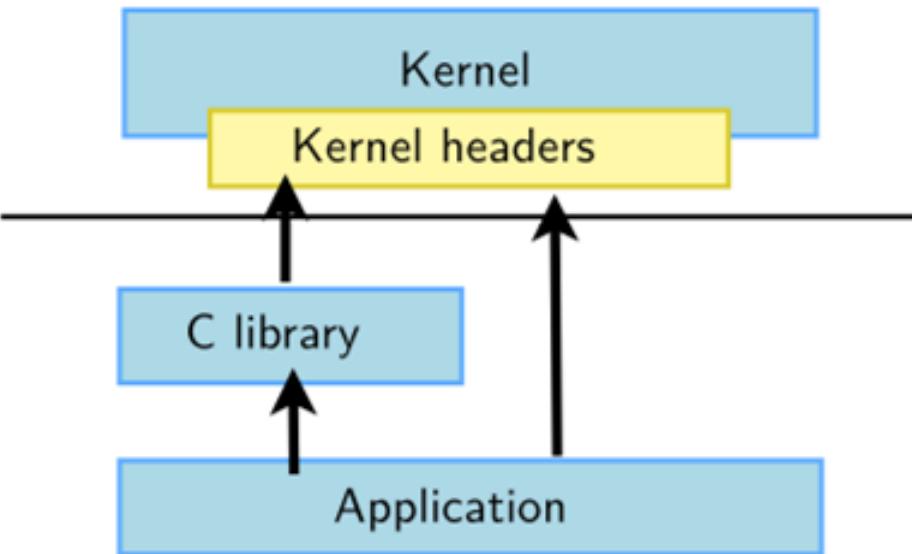
## Kernel Headers - includes



```
include  
acpi  
asm-generic  
clocksource  
crypto  
drm  
dt-bindings  
generated  
keys  
kunit  
kvm  
linux  
math-emu  
media  
memory  
misc  
net  
pcmcia  
ras  
rdma  
rv  
scsi  
soc  
sound  
target  
trace  
uapi  
ufs  
vdso  
video  
xen
```

# DIY LINUX – How? –

## Kernel Headers – architecture dependent



```
bhanu@fedora:~/Documents/tars_glibc/tars/linux/linux-6.7.3$ tree -d -L 2
```

```
.
```

- arch
  - alpha
  - arc
  - arm
  - arm64
  - csky
  - hexagon
  - loongarch
  - m68k
  - microblaze
  - mips
  - nios2
  - openrisc
  - parisc
  - powerpc
  - riscv
  - s390
  - sh
  - sparc
  - um
  - x86
  - xtensa

# DIY LINUX – How? –

## Kernel Headers - drivers

```
drivers
├── accel
├── accessibility
├── acpi
├── amba
├── android
├── ata
├── atm
├── auxdisplay
├── base
├── bcma
├── block
├── bluetooth
├── bus
├── cache
├── cdrom
├── cdx
├── char
├── clk
├── clockssource
├── comedi
├── connector
├── counter
├── cpufreq
├── cpuidle
├── crypto
├── cxl
├── dax
├── dca
├── devfreq
├── dio
├── dma
├── dma-buf
├── dpll
├── edac
├── eisa
├── extcon
├── firewire
├── firmware
├── fpga
└── fsi
```

```
firmware
├── fpga
├── fsi
├── gnss
├── gpio
├── gpu
├── greybus
├── hid
├── hsi
├── hte
├── hv
├── hwmon
├── hwspinlock
├── hwtracing
├── i2c
├── i3c
├── idle
├── iio
├── infiniband
├── input
├── interconnect
├── iommu
├── ipack
├── irqchip
├── isdn
├── leds
├── macintosh
├── mailbox
├── mcb
├── md
├── media
├── memory
├── memstick
├── message
├── mfd
├── misc
├── mmc
├── most
├── mtd
├── mux
├── net
├── nfc
└── ntbus
```

```
net
├── nfc
├── ntb
├── nubus
├── nvdimm
├── nvme
├── nvmem
├── of
├── opp
├── parisc
├── parport
├── pci
├── pcmcia
├── peci
├── perf
├── phy
├── pinctrl
├── platform
├── pmdomain
├── pnp
├── power
├── powercap
├── pps
├── ps3
├── ptp
├── pwm
├── rapido
├── ras
├── regulator
├── remoteproc
├── reset
├── rpmsg
├── rtc
├── s390
├── sbus
├── scsi
├── sh
├── siox
├── slimbus
├── soc
├── soundwire
├── spi
└── spmi
```

```
soundwire
├── spi
├── spmi
├── ssb
├── staging
├── target
├── tc
├── tee
├── thermal
├── thunderbolt
├── tty
├── ufs
├── uio
├── usb
├── vdpa
├── vfio
├── vhost
├── video
├── virt
├── virtio
├── w1
├── watchdog
└── xen
└── zorro
```



# DIY LINUX – How? –

## Kernel Headers – file system

fs	hpfs hugetlbfs iomap isofs jbd2 jffs2 jfs kernfs lockd minix netfs nfs nfs_common nfsd nilfs2 nls notify ntfs ntfs3 ocfs2 omfs openpromfs orangefs overlayfs proc pstore qnx4 qnx6 quota ramfs reiserfs romfs smb squashfs sysfs sysv tracefs ubifs udf ufs unicode vboxsf verity
9p adfs affs afs autofs bcachefs beefs bfs btrfs cachefiles ceph coda configfs cramfs crypto debugfs devpts dlm ecryptfs efivarfs efs erofs exfat exportfs ext2 ext4 f2fs fat freevxfs fscache fuse gfs2 hfs hfsplus hostfs hpfs hugeglbfs iomap	

# DIY LINUX – How?

## Compiler

### GCC Components

CC++ Compiler  
`arm-linux-gnueabihf-gcc`  
`arm-linux-gnueabihf-g++`

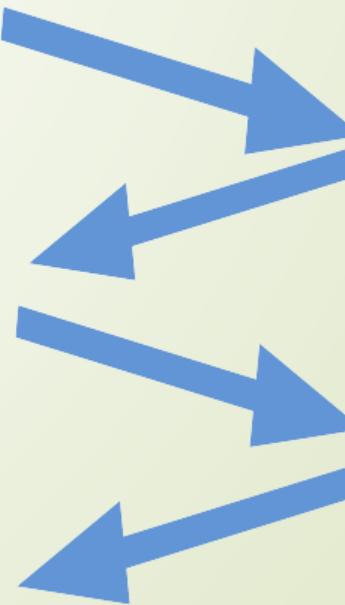
Compiler Support Library  
`libgcc.a / libgcc.so`

Standard C++ Library  
`libstdc++.a / libstdc++.so`

### Glibc Components

Standard C Library Headers  
And Startup Files  
`stdio.h, sthlib.h, pthread.h...`  
`crti.o, crt0.o crt.o/libc.so`

Standard C Library  
`libc.a / libc.so`



# DIY LINUX – How?

## Compiler –Resources to build compiler

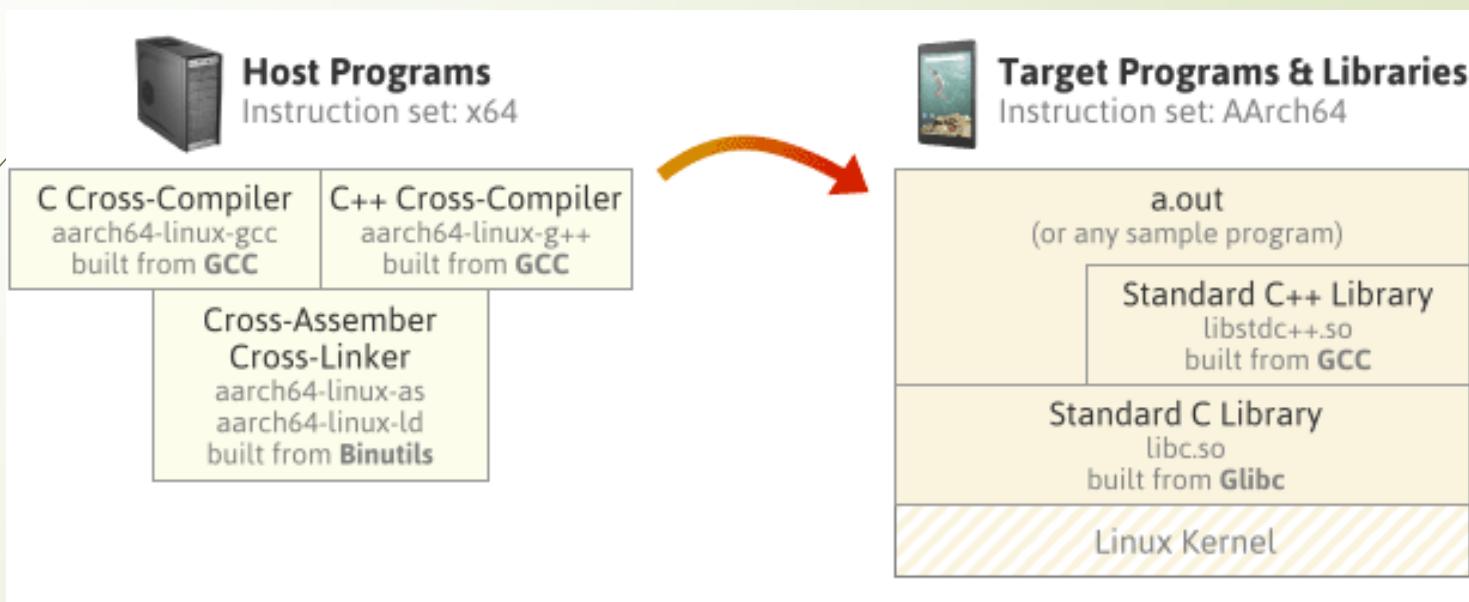
```
bhanu@fedora:~/Documents/tars_glibc/tars/gcc/gcc-13.0.1-20230401$ tree -d -L 1
.
├── config
├── contrib
├── c++tools
├── fixincludes
├── gcc
├── gnatools
├── gotools
└── include
    ├── INSTALL
    └── intl
        ├── libada
        ├── libatomic
        ├── libbacktrace
        ├── libcc1
        ├── libcdy
        ├── libcpp
        ├── libdecnumber
        ├── libffi
        ├── libgcc
        ├── libgfortran
        ├── libgn2
        ├── libgo
        ├── libgomp
        ├── libiberty
        ├── libitm
        ├── libobjc
        ├── libphobos
        ├── libquadmath
        ├── libsanitizer
        ├── libssp
        ├── libstdc++-v3
        ├── libvtv
        └── lto-plugin
            └── maintainer-scripts
```

```
bhanu@fedora:~/Documents/tars_glibc/tars/gcc/gcc-static$ tree -d -L 1
.
├── build-x86_64-cross-linux-gnu
├── fixincludes
├── gcc
├── intl
├── libbacktrace
├── libcdy
├── libcpp
├── libdecnumber
├── libiberty
├── lto-plugin
└── riscv64-unknown-linux-gnu
    └── zlib
```

```
bhanu@fedora:~/Documents/tars_glibc/tars/gcc/gcc-build$ tree -d -L 1
.
├── build-x86_64-cross-linux-gnu
├── build-x86_64-pc-linux-gnu
├── c++tools
├── fixincludes
├── gcc
├── intl
├── libbacktrace
├── libcdy
├── libcpp
├── libdecnumber
├── libiberty
├── lto-plugin
└── x86_64-pc-linux-gnu
    └── zlib
```

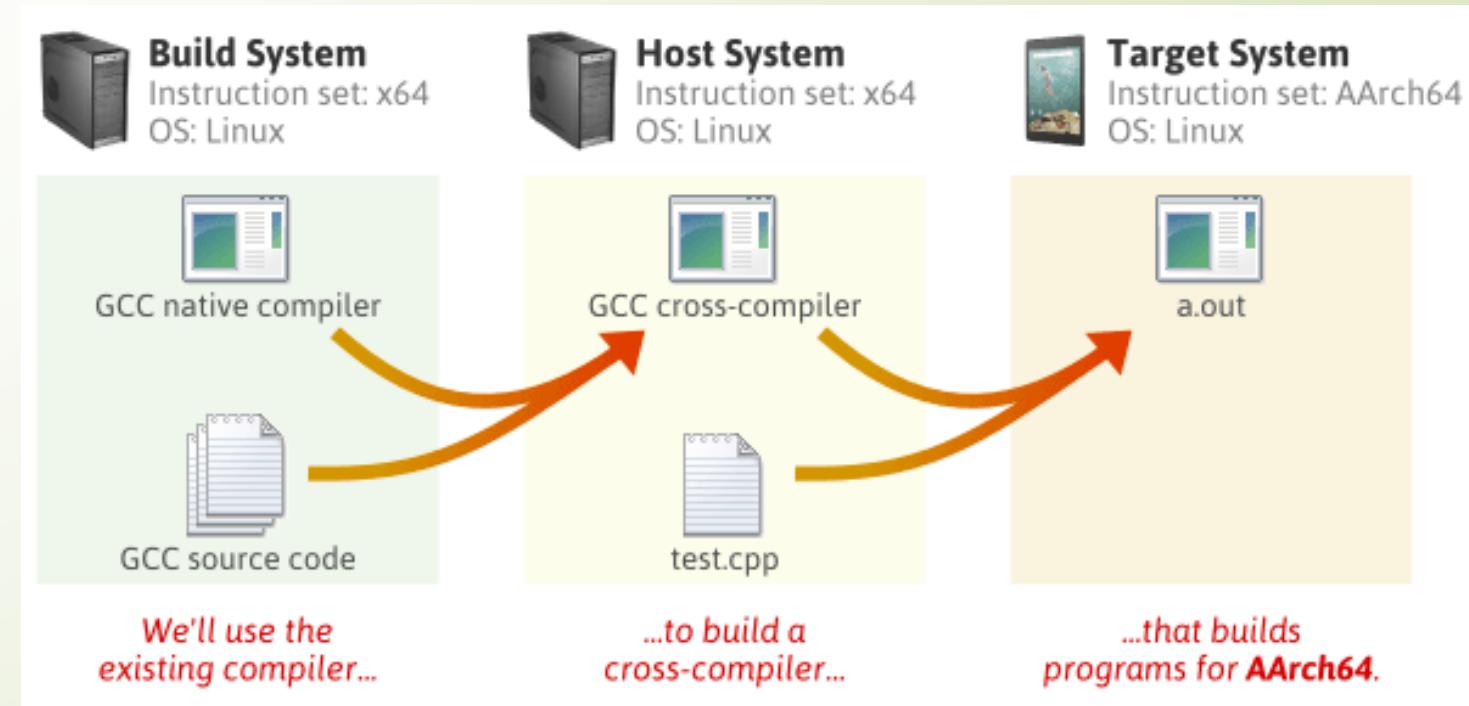
# DIY LINUX – How?

## Cross-Compiler



# DIY LINUX – How?

## Cross-Compiler - build



# DIY LINUX – How?

Cross-Compiler - contains

```
bhanu@fedora:/lj-os/cross-tools$ tree -d -L 2
.
├── bin
├── include
└── lib
    ├── bfd-plugins
    ├── gcc
    ├── lib64 -> lib
    ├── libexec
    └── riscv64-unknown-linux-gnu
        ├── bin
        ├── lib
        ├── share
        ├── info
        └── man
            └── x86_64-pc-linux-gnu
                └── riscv64-unknown-linux-gnu
```

```
bhanu@fedora:/lj-os/cross-tools$ tree -d -L 9
.
├── bin
├── include
└── lib
    ├── bfd-plugins
    └── gcc
        ├── riscv64-unknown-linux-gnu
        │   ├── 13.0.1
        │   │   ├── include
        │   │   ├── include-fixed
        │   │   └── install-tools
        │   └── plugin
        │       ├── ada
        │       ├── gcc-interface
        │       ├── c-family
        │       ├── config
        │       │   └── riscv
        │       ├── cp
        │       ├── d
        │       ├── m2
        │       └── objc
        └── lib64 -> lib
        └── libexec
            └── gcc
                └── riscv64-unknown-linux-gnu
                    ├── 13.0.1
                    │   └── install-tools
                    └── plugin
                    └── riscv64-unknown-linux-gnu
                        ├── bin
                        ├── lib
                        └── ldscripts
                        └── share
                            ├── info
                            └── man
                                ├── man1
                                └── man7
                        └── x86_64-pc-linux-gnu
                            └── riscv64-unknown-linux-gnu
                                ├── include
                                └── lib
```

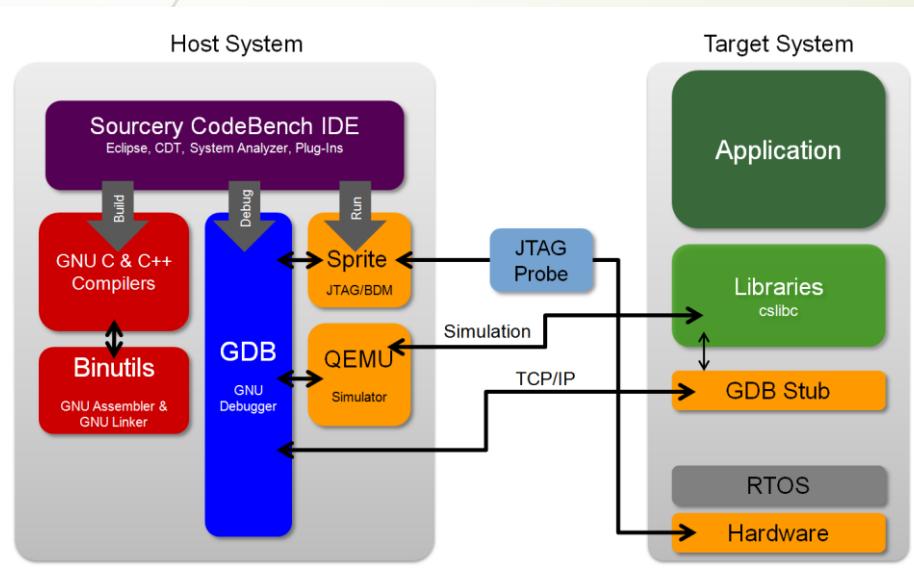
# DIY LINUX – How?

## Cross-Compiler executable test

```
liveuser@localhost-live:~/Documents$ readelf -A test.out
Attribute Section: riscv
File Attributes
  Tag_RISCV_stack_align: 16-bytes
  Tag_RISCV_arch: "rv64i2p1_m2p0_a2p1_f2p2_d2p2_v1p0_zicsr2p0_zifencei2p0_zmmullp0_zve32f1p0_zve32x1p0_zve64d1p0_zve64f1p0_zve64x1p0_zvl128b1p0_zvl32b1p0_zvl64b1p0"
liveuser@localhost-live:~/Documents$
```

# DIY LINUX – How?

## Cross-Compiler - Application





# THANK YOU !